

FULL STACK WEB DEVELOPER

PART XVI

NoSQL Databases

intro to MongoDB



Panos M.

Full Stack Web Developer Part XVI: NoSQL Databases - MongoDB

Panos Matsinopoulos

This book is for sale at

<http://leanpub.com/full-stack-web-developer-part-xvi-no-sql-databases-mongo-db>

This version was published on 2019-08-22



This is a [Leanpub](#) book. Leanpub empowers authors and publishers with the Lean Publishing process. [Lean Publishing](#) is the act of publishing an in-progress ebook using lightweight tools and many iterations to get reader feedback, pivot until you have the right book and build traction once you do.

© 2019 Tech Careet Booster - Panos M.

Contents

The Bundle and the TCB Subscription	1
Part of Bundle	2
Goes with a TCB Subscription	3
Each TCB Subscription Goes with the Bundle	4
Credits To Photo on Cover Page	5
About and Copyright Notice	6
NoSQL Databases	7
1 - Mongo DB	8
Summary	8
Learning Goals	8
Introduction	9
Installing MongoDB	9
Configuration	10
MongoDB Running Process	10
Mongo Client and Mongo Shell	11
Display Current Database	11
Display All Non-Empty Databases	11
Switching Current Database	12
Create a New Database	12
List All Documents	13
CRUD Operations	14
Indexes	24
Ruby Integration	34
Ruby on Rails Integration	39
Closing The MongoDB Chapter	62
Tasks and Quizzes	62

The Bundle and the TCB Subscription

Part of Bundle

This book is not sold alone. It is part of the bundle [Full Stack Web Developer](#).

Goes with a TCB Subscription

When you purchase the bundle, then you have full access to the contents of the [TCB Courses](#).

Each TCB Subscription Goes with the Bundle

Moreover, this goes vice-versa. If you purchase the subscription to the [TCB Courses](#), then you are automatically eligible for the [Full Stack Web Developer](#) bundle.

Credits To Photo on Cover Page

Image by [Michael Gaida](#) from [Pixabay](#)

About and Copyright Notice

Full Stack Web Developer - Part XVI - No SQL Databases

1st Edition, August 2019

by Panos M. for Tech Career Booster (<https://www.techcareerbooster.com>)

Copyright (c) 2019 - Tech Career Booster and Panos M.

All rights reserved. This book may not reproduced in any form, in whole or in part, without written permission from the authors, except in brief quotations in articles or reviews.

Limit of Liability and Disclaimer of Warranty: The author and Tech Career Booster have used their best efforts in preparing this book, and the information provided herein “as is”. The information provided is delivered without warranty, either express or implied. Neither the author nor Tech Career Booster will be held liable for any damages to be caused either directly or indirectly by the contents of the book.

Trademarks: Rather than indicating every occurrence of a trademarked name as such, this book uses the names only in an editorial fashion and to the benefit of the trademark owner with no intention of infringement of the trademark.

For more information: <https://www.techcareerbooster.com>

NoSQL Databases

This book is about MongoDB. It is a very popular example of a NoSQL database. NoSQL databases have become very popular and you will definitely meet them in your software engineering life sooner rather than later. Hence, this introduction is very important so that you will be able to catch up quickly with this new technology.

1 - Mongo DB

Summary



C	• insert()
R	• find()
U	• update()
D	• remove()

mongoDB®

MongoDB - Summary

MongoDB is a very popular document database management system. This chapter gives you enough knowledge to understand what MongoDB is and to be able to start a Web application that is backed up by a MongoDB server for persistency. You will also understand what are the differences to RDBMS systems. We consider this a very important chapter because of the increased popularity of NoSQL databases and MongoDB per se.

Learning Goals

1. Learn what type of database is MongoDB.
2. Learn how to install MongoDB.
3. Learn how to configure MongoDB server.
4. Learn how to check whether MongoDB server is running or not.
5. Learn how to start an interactive MongoDB console.
6. Learn how to display the current database.
7. Learn how to display the non-empty databases.
8. Learn how to switch your current database.
9. Learn how to create a new database.
10. Learn how to create new collections.
11. Learn how to create new documents.
12. Learn how you can list all the documents of a collection.
13. Learn how to create many documents with one command.
14. Learn how to retrieve documents using various criteria.
15. Learn how to use a projection to limit the number of fields returned.
16. Learn how you can limit the number of results returned.
17. Learn how to use special query operators.
18. Learn how to update one or more documents.
19. Learn how to replace one document.
20. Learn about embedded documents.

21. Learn how to query based on fields in embedded documents.
22. Learn about Array values.
23. Learn how to query documents on fields that have Array values.
24. Learn how to iterate over the results using a cursor.
25. Learn how to define indexes on your documents.
26. Learn about the query execution plan.
27. Learn how to list all indexes.
28. Learn how to remove an index.
29. Learn how to create unique indexes.
30. Learn how to create compound indexes.
31. Learn about the way you can integrate Ruby application with MongoDB server.
32. Learn about the way you can integrate Ruby on Rails Web application with MongoDB server.
33. Learn how to use Object to Document mapping gem and create associations between objects.
34. Learn how to add validations to your documents.

Introduction

MongoDB is a *document* database. In other words, in MongoDB a record is a document that resembles JSON objects. For example:

```
1  {
2    first_name: "John",
3    last_name: "Papas",
4    identity_number: "12345678"
5    salary: 3600
6 }
```

(the above code snippet online)

Documents then are characterized by field - value pairs. For example, in the previous document, you can see the field `first_name` with value "John".

Note that the documents can contain other documents. In that case, the document that is contained inside the other is called an embedded document.

Installing MongoDB

You can find installation instructions for MongoDB [here](#). If you have Mac OS X, the instructions are [here](#).

Update your Homebrew

```
1 brew update
```

(the above code snippet online)

Install MongoDB

```
1 brew install mongodb
```

(the above code snippet online)

Add MongoDB Bin Folder To PATH

MongoDB installs shortcuts to its executables inside `/usr/local/bin`. So, if this path is inside your PATH environment variable, then you are good to go. Otherwise, you will have to add it.

Configuration

When MongoDB server is running it is given a path to a configuration file. Usually this is `/usr/local/etc/mongod.conf`. In this file, you specify also the path to the data folder, i.e. the folder in which the documents are going to be saved.

These are the contents of my file:

```
1 systemLog:
2   destination: file
3   path: /usr/local/var/log/mongodb/mongo.log
4   logAppend: true
5 storage:
6   dbPath: /usr/local/var/mongodb
7 net:
8   bindIp: 127.0.0.1
```

(the above code snippet online)

As you can read from above, my data folder is `/usr/local/var/mongodb`.

MongoDB Running Process

You can always check whether MongoDB is running using the following command:

```
1 ps -ef | grep 'mongo'
2   501  834      1  0 12:23PM ??          0:32.64 /usr/local/opt/mongodb/bin/mongo\
3 d --config /usr/local/etc/mongod.conf
```

(the above code snippet online)

The `/usr/local/opt/mongodb/bin/mongod` is the executable that is used to start MongoDB server.

MongoDB is always running and it is a process that listens for client connections. Like MySQL database server. The default port MongoDB server is listening to, is port number 27017.

Mongo Client and Mongo Shell

You can connect to MongoDB server by executing `mongo`. The `mongo` command starts what we call the Mongo shell. It is a JavaScript shell that allows you to execute queries and commands in your MongoDB database. The commands are being sent over to MongoDB server which executes them and returns the results back to the client.

```
1 $ mongo
2 MongoDB shell version v3.4.2
3 connecting to: mongodb://127.0.0.1:27017
4 MongoDB server version: 3.4.2
5 Server has startup warnings:
6 2017-08-08T12:23:05.168+0300 I CONTROL [initandlisten]
7 2017-08-08T12:23:05.168+0300 I CONTROL [initandlisten] ** WARNING: Access contro\|
8 l is not enabled for the database.
9 2017-08-08T12:23:05.168+0300 I CONTROL [initandlisten] ** Read and writ\|
10 e access to data and configuration is unrestricted.
11 2017-08-08T12:23:05.168+0300 I CONTROL [initandlisten]
12 >
```

(the above code snippet online)

The `>` is the mongo shell prompt. Now you can use JavaScript commands, including MongoDB-specific JavaScript commands.

Hint: If you want to exit the mongo shell, then give the command `quit()`. Or hit `<kbd>Ctrl + C</kbd>`.

Display Current Database

Being in a mongo shell session, let's type some commands. The following displays the current database:

```
1 > db
2 test
3 >
```

(the above code snippet online)

As you can see the default database is the one with name `test`.

Display All Non-Empty Databases

```

1 > show dbs
2 blog_development 0.000GB
3 local            0.000GB
4 >

```

(the above code snippet online)

Note that the output of this command, `show dbs`, includes all the non-empty databases. That is why it does not list the `test` database. `test` database is empty.

Note also that the above output will be different to yours.

Switching Current Database

In order to switch to a database, you have to use the helper `use <name_of_database>`. See the following interaction:

```

1 > use local
2 switched to db local
3 > db
4 local
5 >

```

(the above code snippet online)

Create a New Database

In order to create a new database, you have first to `use` it. Then you need to create the first document. Note that documents need to be created inside a collection. Here is an example:

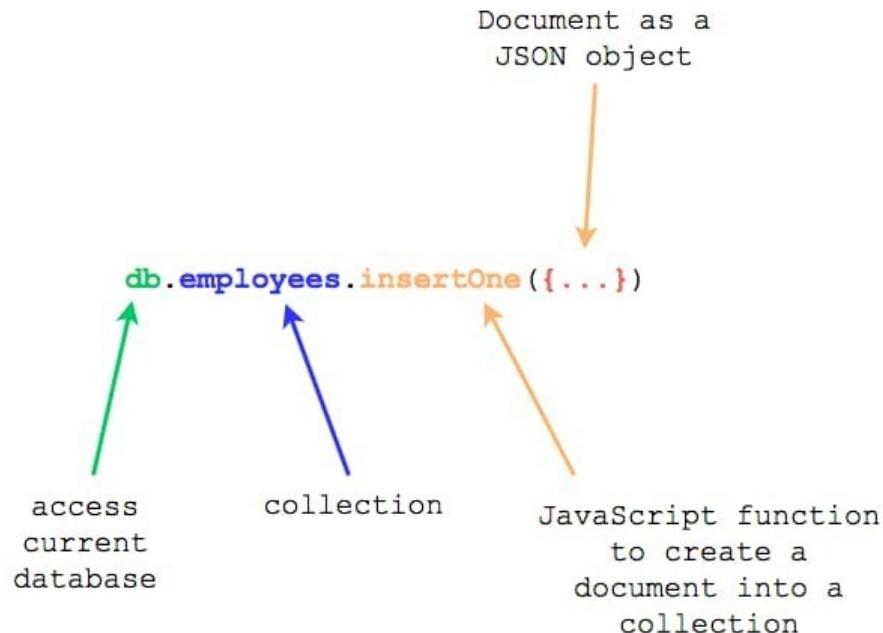
```

1 > use accounting
2 switched to db accounting
3 > db.employees.insertOne({first_name: "John", last_name: "Papas", identity_number\
4 : "12345678", salary: 3600})
5 {
6     "acknowledged" : true,
7     "insertedId" : ObjectId("598a10b21631fe0993ee7974")
8 }
9 > show dbs
10 accounting      0.000GB
11 blog_development 0.000GB
12 local           0.000GB
13 >

```

(the above code snippet online)

1. Line 1 switches to a new database.
2. Line 3 creates a **new document** within a **new collection** within a **new database**
3. Line 8 proves now that accounting database is one of the non-empty databases.



How to Create Database alongside the First Collection and the First Document

In other words, in order to create a new database, you need to create a new collection and a new document inside that new collection. MongoDB will receive the command `db.<collection>.insertOne(<document>)` and will create the collection and the database if they don't exist, before actually creating the document itself.

List All Documents

You can list the documents of a collection as follows:

```

1 > db.employees.find();
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "John", "last_name" :
3 " " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 >

```

(the above code snippet online)

The above will return all the documents of the collection in groups of 20. Do you see the "_id" property? This is automatically assigned by MongoDB in order to make the documents unique. It works like the primary key of the document.

Another way to list all documents is using the .pretty() function:

```

1 > db.employees.find().pretty();
2 {
3     "_id" : ObjectId("598a9d09c0317e30b80c33a5"),
4     "first_name" : "John",
5     "last_name" : "Papas",
6     "identity_number" : "12345678",
7     "salary" : 3600
8 }
9 >

```

(the above code snippet online)

CRUD Operations

Let's see some basics with regards to CRUD (Create, Retrieve, Update, Delete) using mongo shell.

Create

In order to create a new document you need to use either `insertOne()` or `insertMany()` methods on the collection reference.

Let's create some documents inside the `employees` collection (in our `accounting` db).

```

1 > db.employees.insertOne({first_name: "Mary", last_name: "Foo", identity_number: \
2 "87654321", "salary": 2800});
3 {
4     "acknowledged" : true,
5     "insertedId" : ObjectId("598ab700c0317e30b80c33a6")
6 }
7 >

```

(the above code snippet online)

The above creates a new document. Let's create more than one:

```

1 > db.employees.insertMany([
2 ... {first_name: "Paul", last_name: "Bar", identity_number: "37281493", salary: 3\
3 500},
4 ... {first_name: "Peter", last_name: "Pan", identity_number: "38589242123", salar\
5 y: 1500}
6 ... ]);
7 {
8     "acknowledged" : true,
9     "insertedIds" : [
10         ObjectId("598ab79ec0317e30b80c33a7"),
11         ObjectId("598ab79ec0317e30b80c33a8")
12     ]
13 }
14 >

```

(the above code snippet online)

The `insertMany()` above creates two documents to the `employees` collection. Let's now get the list of all the documents in the collection:

```

1 > db.employees.find()
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "John", "last_name\"\
3 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\"\
5 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
6 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "first_name" : "Paul", "last_name\"\
7 " : "Bar", "identity_number" : "37281493", "salary" : 3500 }
8 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\
9 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
10 >

```

(the above code snippet online)

Retrieve Documents

The `find()` is the main method that is used to retrieve documents from the database.

If you don't give any arguments it will return all the documents. Otherwise, the first argument is the filtering/where clause.

```

1 > db.employees.find({first_name: "Paul"})
2 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "first_name" : "Paul", "last_name\"\
3 " : "Bar", "identity_number" : "37281493", "salary" : 3500 }
4 >

```

(the above code snippet online)

The above query brings all the employees that have `first_name` equal to "Paul".

You can also give a second argument to the `find()` method. The second argument works like a projection/fields selection argument.

```

1 > db.employees.find({first_name: "Paul"}, {first_name: 1, identity_number: 1})
2 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "first_name" : "Paul", "identity_\n3 number" : "37281493" }
4 >

```

(the above code snippet online)

The above projection selects the properties `first_name` and `identity_number`. Please, note that the `_id` property is always returned.

```

1 > db.employees.find({first_name: "Paul"}, {first_name: 0})
2 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "last_name" : "Bar", "identity_nu\
3 mber" : "37281493", "salary" : 3500 }
4 >

```

(the above code snippet online)

The `{first_name: 0}`, on the other hand, specifies which property should **not** be returned.

Note that the result of the `find()` is a *cursor*. It is the mongo shell that uses the cursor to iterate and scan through the results.

You can also modify the cursor, by specifying things like the `limit`, for example:

```

1 > db.employees.find().limit(2)
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "John", "last_name\
3 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\
5 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
6 >

```

(the above code snippet online)

The above limits the cursor results to the first 2 records located.

Your filter may contain query operators. For example:

```

1 > db.employees.find({salary: {$gt: 2500}})
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name\
3 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\
5 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
6 >

```

(the above code snippet online)

The `{salary: {$gt: 2500}}` filter is using the query operator `$gt`, i.e. greater than. There are other query operators that you can use. You can read more about them [here](#).

Another more complicated example is this one:

```

1 > db.employees.find({$or: [{salary: {$gt: 2900}}, {first_name: "Peter"}]}));
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name\"
3 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\"
5 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
6 >

```

(the above code snippet online)

The `[$or: [{salary: {$gt: 2900}}, {first_name: "Peter"}]]` is using the `$or` operator. The value of `$or` is an array of conditions that will be ORed. If any of the conditions matches, then document will be included in the result returned.

The [MongoDB documentation page on querying documents](#) is a very good reference to visit every time you want to create a query and you have doubts about.

Querying using JavaScript becomes more complicated when the document has a more complicated structure.

We will see more advanced queries later on.

Update Documents

The methods to use to update existing documents are `updateOne()`, `updateMany()` and `replaceOne()`.

Let's suppose that we want to change the name of the employee with `first_name` "John" to be "Nick":

```

1 > db.employees.find()
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "John", "last_name\"
3 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\"
5 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
6 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "first_name" : "Paul", "last_name\"
7 " : "Bar", "identity_number" : "37281493", "salary" : 3500 }
8 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\"
9 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
10 > db.employees.updateOne({first_name: "John"}, {$set: {first_name: "Nick"}});
11 { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
12 > db.employees.find()
13 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name\"
14 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
15 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\"
16 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
17 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "first_name" : "Paul", "last_name\"
18 " : "Bar", "identity_number" : "37281493", "salary" : 3500 }
19 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\"
20 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
21 >

```

(the above code snippet online)

What happens above? We first display all the employees. Then we use `updateOne()` to specify which document to update and how.

The `updateOne()` takes as first argument the *filter/where* object. It tells which document is going to be updated. In the above example we specify that the document to be updated is the one with `first_name` equal to "John". The second argument to `updateOne()` is the update action. We use `$set` to tell that we want to set specific fields. In the example above we update the `first_name`.

The `updateMany()` is similar to `updateOne()` with regards to its syntax. But the `updateMany()` is updating all the matching documents, whereas the `updateOne()` updates only one document even if more than one might match the filter.

The `replaceOne()` method is used to replace a document with another one:

```

1 > db.employees.replaceOne({first_name: "Paul"}, {first_name: "Keith", last_name: \
2 "Reilly", identity_number: "3728817", salary: 2900})
3 { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
4 > db.employees.find()
5 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name" \
6 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
7 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name" \
8 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
9 { "_id" : ObjectId("598ab79ec0317e30b80c33a7"), "first_name" : "Keith", "last_name" \
10 " : "Reilly", "identity_number" : "3728817", "salary" : 2900 }
11 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_name" \
12 " : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
13 >

```

(the above code snippet online)

The above command replaces the document that has `first_name` equal to "Paul", with a completely new document. Note, however, that the command does not replace the id of the document. Hence, the new document has the same id as the previous one.

Delete Documents

The methods to delete documents from a collection are `deleteOne()` and `deleteMany()`.

The first one, matches a single document using its first argument as a filter/where clause. The second one will match many documents.

```

1 > db.employees.deleteOne({first_name: "Keith"})
2 { "acknowledged" : true, "deletedCount" : 1 }
3 > db.employees.find()
4 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name\\
5 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
6 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\\
7 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
8 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\\
9 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
10 >

```

(the above code snippet online)

Embedded/Nested Documents

Let's create some more complicated documents. Assume that we want to add `contact_details` on each one of our employees:

```

1 > db.employees.find()
2 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name\\
3 " : "Papas", "identity_number" : "12345678", "salary" : 3600 }
4 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\\
5 " : "Foo", "identity_number" : "87654321", "salary" : 2800 }
6 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\\
7 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500 }
8 > db.employees.updateOne({first_name: "Nick"}, {$set: {contact_details: {address:\\
9 "1A Meyrick Park Cres", city: "Bournemouth", zip_code: "BH3 7AG", country: "UK"}\\
10 }})
11 { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
12 > db.employees.updateOne({first_name: "Mary"}, {$set: {contact_details: {address:\\
13 "6 Coronach Way", city: "New Rossington", zip_code: "DN11 0RN", country: "UK"}}}\\
14 )
15 { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
16 > db.employees.updateOne({first_name: "Peter"}, {$set: {contact_details: {address:\\
17 "38 Coleraine Rd", city: "London", zip_code: "N8 0QL", country: "UK"}}})
18 { "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
19 > db.employees.find()
20 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name\\
21 " : "Papas", "identity_number" : "12345678", "salary" : 3600, "contact_details" :\\
22 { "address" : "1A Meyrick Park Cres", "city" : "Bournemouth", "zip_code" : "BH3 \\
23 7AG", "country" : "UK" } }
24 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name\\
25 " : "Foo", "identity_number" : "87654321", "salary" : 2800, "contact_details" : {\\
26 "address" : "6 Coronach Way", "city" : "New Rossington", "zip_code" : "DN11 0RN"\\
27 , "country" : "UK" } }
28 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_nam\\

```

```

29 e" : "Pan", "identity_number" : "38589242123", "salary" : 1500, "contact_details"\ 
30 : { "address" : "38 Coleraine Rd", "city" : "London", "zip_code" : "N8 0QL", "co\ 
31 untry" : "UK" } } 
32 >

```

(the above code snippet online)

We have used the `updateOne()` method to add an embedded/nested document inside each one of the existing documents. Now all the employees have a nested document as value of their `contact_details` property.

You can use the dot notation to refer to fields in the embedded documents. For example:

```

1 > db.employees.find({ "contact_details.city" : "London" })
2 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_name" : "Smith", "email" : "peter.smith@example.com", "contact_details" : { "address" : "38 Coleraine Rd", "city" : "London", "zip_code" : "N8 0QL", "country" : "UK" } }
3 >

```

(the above code snippet online)

The above query returns all the employees that have contact details city equal to "London".

Array Values

Another case that it is interesting, is the case in which some of the document properties have as value an Array.

Let's do an example. We are going to use another database, named `book_store`. The following commands create the collection `books`.

```

1 > use book_store;
2 switched to db book_store
3 > db.books.insertOne({title: "Of Mice And Men", author: "John Steinbeck", tags: ["novel", "classic"]});
4 {
5     "acknowledged" : true,
6     "insertedId" : ObjectId("598b5ee4c0317e30b80c33a9")
7 }
8 > db.books.insertOne({title: "The Idiot", author: "Fyodor Dostoyevsky", tags: ["novel", "classic"]});
9 {
10    "acknowledged" : true,
11    "insertedId" : ObjectId("598b5f20c0317e30b80c33aa")
12 }
13 > db.books.insertOne({title: "Life of Pi", author: "Yann Martel", tags: ["novel", "modern"]});
14

```

```

17 {
18     "acknowledged" : true,
19     "insertedId" : ObjectId("598b5f5bc0317e30b80c33ab")
20 }
21 > db.books.insertOne({title: "Protable Cosmos", author: "Jones Alexander", tags: \
22 ["history", "modern"]});
23 {
24     "acknowledged" : true,
25     "insertedId" : ObjectId("598b5fa5c0317e30b80c33ac")
26 }
27 > db.books.find()
28 { "_id" : ObjectId("598b5ee4c0317e30b80c33a9"), "title" : "Of Mice And Men", "aut\
29 hor" : "John Steinbeck", "tags" : [ "novel", "classic" ] }
30 { "_id" : ObjectId("598b5f20c0317e30b80c33aa"), "title" : "The Idiot", "author" :\
31 "Fyodor Dostoyevsky", "tags" : [ "novel", "classic" ] }
32 { "_id" : ObjectId("598b5f5bc0317e30b80c33ab"), "title" : "Life of Pi", "author" \
33 : "Yann Martel", "tags" : [ "novel", "modern" ] }
34 { "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title" : "Protable Cosmos", "aut\
35 hor" : "Jones Alexander", "tags" : [ "history", "modern" ] }
36 >

```

(the above code snippet online)

Each document has a property named tags which has an Array value. Let's do some queries:

```

1 > db.books.find({tags: ["history", "modern"]});
2 { "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title" : "Protable Cosmos", "aut\
3 hor" : "Jones Alexander", "tags" : [ "history", "modern" ] }
4 >

```

(the above code snippet online)

The filter {tags: ["history", "modern"]} tells MongoDB to return books that they have tags property equal to ["history", "modern"].

```

1 > db.books.find({tags: ["modern", "history"]});
2 >

```

(the above code snippet online)

Note the above. If you try the same tag but with different order of the Array values, the document will not match. It is a different query. The match of the Array values needs to be exact and with the same order.

If you don't care about the order of the values, then you need to use the \$all query operator:

```

1 > db.books.find({tags: {$all: ["modern", "history"]}});
2 { "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title" : "Protable Cosmos", "aut\
3 hor" : "Jones Alexander", "tags" : [ "history", "modern" ] }
4 >

```

(the above code snippet online)

If you want to find documents that have a specific value into their tags Array, the query becomes much simpler:

```

1 > db.books.find({tags: "novel"})
2 { "_id" : ObjectId("598b5ee4c0317e30b80c33a9"), "title" : "Of Mice And Men", "aut\
3 hor" : "John Steinbeck", "tags" : [ "novel", "classic" ] }
4 { "_id" : ObjectId("598b5f20c0317e30b80c33aa"), "title" : "The Idiot", "author" :\\
5 "Fyodor Dostoyevsky", "tags" : [ "novel", "classic" ] }
6 { "_id" : ObjectId("598b5f5bc0317e30b80c33ab"), "title" : "Life of Pi", "author" \
7 : "Yann Martel", "tags" : [ "novel", "modern" ] }
8 >

```

(the above code snippet online)

The following query brings the books that have one of the tags "classic" or "history":

```

1 > db.books.find({$or: [{tags: "classic"}, {tags: "history"}]})
2 { "_id" : ObjectId("598b5ee4c0317e30b80c33a9"), "title" : "Of Mice And Men", "aut\
3 hor" : "John Steinbeck", "tags" : [ "novel", "classic" ] }
4 { "_id" : ObjectId("598b5f20c0317e30b80c33aa"), "title" : "The Idiot", "author" :\\
5 "Fyodor Dostoyevsky", "tags" : [ "novel", "classic" ] }
6 { "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title" : "Protable Cosmos", "aut\
7 hor" : "Jones Alexander", "tags" : [ "history", "modern" ] }
8 >

```

(the above code snippet online)

Iterating Using a Cursor

You can save the query cursor into a variable and then iterate over the results one by one:

```

1  > var books = db.books.find()
2  > while (books.hasNext()) {
3  ... print(tojson(books.next()));
4  ...
5  {
6      "_id" : ObjectId("598b5ee4c0317e30b80c33a9"),
7      "title" : "Of Mice And Men",
8      "author" : "John Steinbeck",
9      "tags" : [
10          "novel",
11          "classic"
12      ]
13 }
14 {
15     "_id" : ObjectId("598b5f20c0317e30b80c33aa"),
16     "title" : "The Idiot",
17     "author" : "Fyodor Dostoyevsky",
18     "tags" : [
19         "novel",
20         "classic"
21     ]
22 }
23 {
24     "_id" : ObjectId("598b5f5bc0317e30b80c33ab"),
25     "title" : "Life of Pi",
26     "author" : "Yann Martel",
27     "tags" : [
28         "novel",
29         "modern"
30     ]
31 }
32 {
33     "_id" : ObjectId("598b5fa5c0317e30b80c33ac"),
34     "title" : "Protable Cosmos",
35     "author" : "Jones Alexander",
36     "tags" : [
37         "history",
38         "modern"
39     ]
40 }
41 >

```

(the above code snippet online)

The method `hasNext()` returns `true` if there are more data to fetch. And the `next()` returns the actual document. Note that we use the `print()` and `tojson()` helpers to print out to the console.

Indexes

Indexes help you find documents quickly, like the indexes we use in RDBMS databases.

Unique Index on `_id` field

As we said earlier, the `_id` field is present in all documents. And it is the primary key. Uniquely identifies a document and two documents cannot have the same id within the same collection.

```

1  > db.books.find()
2  { "_id" : ObjectId("598b5ee4c0317e30b80c33a9"), "title" : "Of Mice And Men", "aut\
3  hor" : "John Steinbeck", "tags" : [ "novel", "classic" ] }
4  { "_id" : ObjectId("598b5f20c0317e30b80c33aa"), "title" : "The Idiot", "author" :\ \
5  "Fyodor Dostoyevsky", "tags" : [ "novel", "classic" ] }
6  { "_id" : ObjectId("598b5f5bc0317e30b80c33ab"), "title" : "Life of Pi", "author" \ \
7  : "Yann Martel", "tags" : [ "novel", "modern" ] }
8  { "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title" : "Protable Cosmos", "aut\
9  hor" : "Jones Alexander", "tags" : [ "history", "modern" ] }
10 >
11 > db.books.insertOne({ "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title": "dum\
12 my", "author" : "dummy", "tags" : [] })
13 2017-08-10T08:02:03.603+0300 E QUERY      [thread1] WriteError: E11000 duplicate ke\
14 y error collection: book_store.books index: _id_ dup key: { : ObjectId('598b5fa5c\
15 0317e30b80c33ac') } :
16 WriteError({
17     "index" : 0,
18     "code" : 11000,
19     "errmsg" : "E11000 duplicate key error collection: book_store.books index\
20 : _id_ dup key: { : ObjectId('598b5fa5c0317e30b80c33ac') }",
21     "op" : {
22         "_id" : ObjectId("598b5fa5c0317e30b80c33ac"),
23         "title" : "dummy",
24         "author" : "dummy",
25         "tags" : [ ]
26     }
27 })
28 WriteError@src/mongo/shell/bulk_api.js:469:48
29 Bulk/mergeBatchResults@src/mongo/shell/bulk_api.js:836:49
30 Bulk/executeBatch@src/mongo/shell/bulk_api.js:906:13
31 Bulk/this.execute@src/mongo/shell/bulk_api.js:1150:21
32 DBCollection.prototype.insertOne@src/mongo/shell/crud_api.js:242:9
33 @(shell):1:1
34 >
```

(the above code snippet online)

The `db.books.insertOne()` command above, failed because we tried to create a new document in the book collection using an `_id` value that already exists for another document. Do you see the error thrown? `E11000 duplicate key error`

Create Index

We can create an index on a single document field. These are called single field indexes. Look at the following example:

```

1 > db.books.createIndex({title: 1})
2 {
3     "createdCollectionAutomatically" : false,
4     "numIndexesBefore" : 1,
5     "numIndexesAfter" : 2,
6     "ok" : 1
7 }
8 >

```

(the above code snippet online)

We have created an index on the field `title` of the `books` collection. The `1` given as a value to the `title` field index specification means that we want the index to be sorting elements in ascending order. If we gave `-1` instead, then we would have specified the index to have the elements ordered in descending order. However, for single field indexes, this is irrelevant, because MongoDB can traverse the index in both directions with the same performance.

Here is a query that uses the `title` as a filter:

```

1 > db.books.find({title: "Life of Pi"})
2 { "_id" : ObjectId("598b5f5bc0317e30b80c33ab"), "title" : "Life of Pi", "author" \
3 : "Yann Martel", "tags" : [ "novel", "modern" ] }
4 >

```

(the above code snippet online)

Execution Plan

The above query, with the `title` index in place, uses the index to quickly fetch the document requested and does not do full collection scan. You can confirm that by asking the execution plan of the query:

```

1 > db.books.find({title: "Life of Pi"}).explain("executionStats")
2 {
3     "queryPlanner" : {
4         "plannerVersion" : 1,
5         "namespace" : "book_store.books",
6         "indexFilterSet" : false,
7         "parsedQuery" : {
8             "title" : {
9                 "$eq" : "Life of Pi"
10            }
11        },
12        "winningPlan" : {
13            "stage" : "FETCH",
14            "inputStage" : {
15                "stage" : "IXSCAN",
16                "keyPattern" : {
17                    "title" : 1
18                },
19                "indexName" : "title_1",
20                "isMultiKey" : false,
21                "multiKeyPaths" : {
22                    "title" : [ ]
23                },
24                "isUnique" : false,
25                "isSparse" : false,
26                "isPartial" : false,
27                "indexVersion" : 1,
28                "direction" : "forward",
29                "indexBounds" : {
30                    "title" : [
31                        "[\"Life of Pi\", \"Life of Pi\"]"
32                    ]
33                }
34            }
35        },
36        "rejectedPlans" : [ ]
37    },
38    "executionStats" : {
39        "executionSuccess" : true,
40        "nReturned" : 1,
41        "executionTimeMillis" : 24,
42        "totalKeysExamined" : 1,
43        "totalDocsExamined" : 1,
44        "executionStages" : {
45            "stage" : "FETCH",
46            "nReturned" : 1,

```

```
47      "executionTimeMillisEstimate" : 0,
48      "works" : 2,
49      "advanced" : 1,
50      "needTime" : 0,
51      "needYield" : 0,
52      "saveState" : 1,
53      "restoreState" : 1,
54      "isEOF" : 1,
55      "invalidates" : 0,
56      "docsExamined" : 1,
57      "alreadyHasObj" : 0,
58      "inputStage" : {
59          "stage" : "IXSCAN",
60          "nReturned" : 1,
61          "executionTimeMillisEstimate" : 0,
62          "works" : 2,
63          "advanced" : 1,
64          "needTime" : 0,
65          "needYield" : 0,
66          "saveState" : 1,
67          "restoreState" : 1,
68          "isEOF" : 1,
69          "invalidates" : 0,
70          "keyPattern" : {
71              "title" : 1
72          },
73          "indexName" : "title_1",
74          "isMultiKey" : false,
75          "multiKeyPaths" : {
76              "title" : [ ]
77          },
78          "isUnique" : false,
79          "isSparse" : false,
80          "isPartial" : false,
81          "indexVersion" : 1,
82          "direction" : "forward",
83          "indexBounds" : {
84              "title" : [
85                  "[\"Life of Pi\", \"Life of Pi\"]"
86              ]
87          },
88          "keysExamined" : 1,
89          "seeks" : 1,
90          "dupsTested" : 0,
91          "dupsDropped" : 0,
92          "seenInvalidated" : 0
```

```

93         }
94     }
95 },
96 "serverInfo" : {
97     "host" : "Panayotiss-MacBook-Pro.local",
98     "port" : 27017,
99     "version" : "3.4.6",
100    "gitVersion" : "c55eb86ef46ee7aede3b1e2a5d184a7df4bfb5b5"
101   },
102  "ok" : 1
103 }
104 >

```

(the above code snippet online)

The `.explain("executionStats")` method returns back information about the method used to fetch the document, execute the query. In the example above, you can see that it returned IXSCAN for the `inputStage.stage` value, which means it used the index. You can also see that the `totalDocsExamined` value is 1, even if the documents / books are 4.

On the other hand, if you try searching by `author` that does not have an index, you will see a full collection scan query explanation:

```

1 > db.books.find({author: "Yann Martel"}).explain("executionStats")
2 {
3     "queryPlanner" : {
4         "plannerVersion" : 1,
5         "namespace" : "book_store.books",
6         "indexFilterSet" : false,
7         "parsedQuery" : {
8             "author" : {
9                 "$eq" : "Yann Martel"
10            }
11        },
12        "winningPlan" : {
13            "stage" : "COLLSCAN",
14            "filter" : {
15                "author" : {
16                    "$eq" : "Yann Martel"
17                }
18            },
19            "direction" : "forward"
20        },
21        "rejectedPlans" : [ ]
22    },
23    "executionStats" : {
24        "executionSuccess" : true,

```

```

25      "nReturned" : 1,
26      "executionTimeMillis" : 0,
27      "totalKeysExamined" : 0,
28      "totalDocsExamined" : 4,
29      "executionStages" : {
30          "stage" : "COLLSCAN",
31          "filter" : {
32              "author" : {
33                  "$eq" : "Yann Martel"
34              }
35          },
36          "nReturned" : 1,
37          "executionTimeMillisEstimate" : 0,
38          "works" : 6,
39          "advanced" : 1,
40          "needTime" : 4,
41          "needYield" : 0,
42          "saveState" : 0,
43          "restoreState" : 0,
44          "isEOF" : 1,
45          "invalidates" : 0,
46          "direction" : "forward",
47          "docsExamined" : 4
48      }
49  },
50  "serverInfo" : {
51      "host" : "Panayotiss-MacBook-Pro.local",
52      "port" : 27017,
53      "version" : "3.4.6",
54      "gitVersion" : "c55eb86ef46ee7aede3b1e2a5d184a7df4bfb5b5"
55  },
56  "ok" : 1
57 }
58 >

```

(the above code snippet online)

Do you see the `inputStage.stage` that has the value `COLLSCAN`? Also do you see the `totalDocsExamined`? It has the value 4. These are indicators that we had a full scan collection and no index has been used.

List All Indexes

You can list all indexes of a collection with the method call `.getIndexes()`:

```

1 > db.books.getIndexes()
2 [
3     {
4         "v" : 1,
5         "key" : {
6             "_id" : 1
7         },
8         "name" : "_id_",
9         "ns" : "book_store.books"
10    },
11    {
12        "v" : 1,
13        "key" : {
14            "title" : 1
15        },
16        "name" : "title_1",
17        "ns" : "book_store.books"
18    }
19 ]
20 >

```

(the above code snippet online)

As you can see above, we have 2 indexes in the books collection. One on _id field and another one on title field:

Drop Index

In order to remove an index, you have to use the dropIndex() method. Let's remove the title index:

```

1 > db.books.dropIndex({title: 1})
2 { "nIndexesWas" : 2, "ok" : 1 }
3 > db.books.getIndexes()
4 [
5     {
6         "v" : 1,
7         "key" : {
8             "_id" : 1
9         },
10        "name" : "_id_",
11        "ns" : "book_store.books"
12    }
13 ]
14 >

```

(the above code snippet online)

The above removed the `{title: 1}` index. As you can see, then call to `getIndexes()` now returns one index only.

Unique Index

Like in RDBM systems, you can create a *unique index* that will protect you from adding documents having the same value on a particular field.

Let's create a unique index on the `title` field:

```

1 > db.books.createIndex({title: 1}, {unique: true})
2 {
3     "createdCollectionAutomatically" : false,
4     "numIndexesBefore" : 1,
5     "numIndexesAfter" : 2,
6     "ok" : 1
7 }
8 >

```

(the above code snippet online)

Now if you try to create a document that has a `title` equal to the `title` of an existing document, then your operation will fail.

```

1 > db.books.find()
2 { "_id" : ObjectId("598b5ee4c0317e30b80c33a9"), "title" : "Of Mice And Men", "aut\
3 hor" : "John Steinbeck", "tags" : [ "novel", "classic" ] }
4 { "_id" : ObjectId("598b5f20c0317e30b80c33aa"), "title" : "The Idiot", "author" :\ \
5 "Fyodor Dostoyevsky", "tags" : [ "novel", "classic" ] }
6 { "_id" : ObjectId("598b5f5bc0317e30b80c33ab"), "title" : "Life of Pi", "author" \
7 : "Yann Martel", "tags" : [ "novel", "modern" ] }
8 { "_id" : ObjectId("598b5fa5c0317e30b80c33ac"), "title" : "Protable Cosmos", "aut\
9 hor" : "Jones Alexander", "tags" : [ "history", "modern" ] }
10 >
11 > db.books.insertOne({title: "Life of Pi", "author": "Unknown", "tags": []})
12 2017-08-10T16:29:06.021+0300 E QUERY      [thread1] WriteError: E11000 duplicate ke\
13 y error collection: book_store.books index: title_1 dup key: { : "Life of Pi" } :
14 WriteError{
15     "index" : 0,
16     "code" : 11000,
17     "errmsg" : "E11000 duplicate key error collection: book_store.books index\
18 : title_1 dup key: { : \"Life of Pi\" }",
19     "op" : {
20         "_id" : ObjectId("598c5fa2c0317e30b80c33ad"),
21         "title" : "Life of Pi",
22         "author" : "Unknown",
23         "tags" : [ ]

```

```

24         }
25     })
26 WriteError@src/mongo/shell/bulk_api.js:469:48
27 Bulk/mergeBatchResults@src/mongo/shell/bulk_api.js:836:49
28 Bulk/executeBatch@src/mongo/shell/bulk_api.js:906:13
29 Bulk/this.execute@src/mongo/shell/bulk_api.js:1150:21
30 DBCollection.prototype.insertOne@src/mongo/shell/crud_api.js:242:9
31 @(shell):1:1
32 >

```

(the above code snippet online)

Do you see the error E11000 duplicate key error collection when we try to create another book instance with the title Life Of Pi? This error is raised thanks to the unique index on the title field.

Compound Index

Besides the single field indexes, one can create compound indexes, i.e. indexes that index the combined values of two or more fields of the collection.

Let's switch to our accounting database and create a unique compound index on the employees collection as follows:

```

1 > use accounting
2 switched to db accounting
3 >
4 > db.employees.find()
5 { "_id" : ObjectId("598a9d09c0317e30b80c33a5"), "first_name" : "Nick", "last_name" :
6 " : "Papas", "identity_number" : "12345678", "salary" : 3600, "contact_details" : \
7 { "address" : "1A Meyrick Park Cres", "city" : "Bournemouth", "zip_code" : "BH3 \
8 7AG", "country" : "UK" } }
9 { "_id" : ObjectId("598ab700c0317e30b80c33a6"), "first_name" : "Mary", "last_name" :
10 " : "Foo", "identity_number" : "87654321", "salary" : 2800, "contact_details" : { \
11 "address" : "6 Coronach Way", "city" : "New Rossington", "zip_code" : "DN11 0RN" \
12 , "country" : "UK" } }
13 { "_id" : ObjectId("598ab79ec0317e30b80c33a8"), "first_name" : "Peter", "last_name" :
14 " : "Pan", "identity_number" : "38589242123", "salary" : 1500, "contact_details" \
15 : { "address" : "38 Coleraine Rd", "city" : "London", "zip_code" : "N8 0QL", "co \
16 untry" : "UK" } }
17 >
18 > db.employees.createIndex({first_name: 1, last_name: 1}, {unique: true})
19 {
20     "createdCollectionAutomatically" : false,
21     "numIndexesBefore" : 1,
22     "numIndexesAfter" : 2,
23     "ok" : 1

```

```

24 }
25 >
26 > db.employees.getIndexes()
27 [
28     {
29         "v" : 1,
30         "key" : {
31             "_id" : 1
32         },
33         "name" : "_id_",
34         "ns" : "accounting.employees"
35     },
36     {
37         "v" : 1,
38         "unique" : true,
39         "key" : {
40             "first_name" : 1,
41             "last_name" : 1
42         },
43         "name" : "first_name_1_last_name_1",
44         "ns" : "accounting.employees"
45     }
46 ]
47 >

```

(the above code snippet online)

The call `db.employees.createIndex({first_name: 1, last_name: 1}, {unique: true})` creates a unique index on the combination of `first_name` and `last_name`.

This allows us to create a document with `first_name` equal to `Peter` like the following:

```

1 > db.employees.insertOne({first_name: "Peter", last_name: "Gorilla", identity_num\
2 ber: "37281923", salary: 3500, contact_details: {}})
3 {
4     "acknowledged" : true,
5     "insertedId" : ObjectId("598c6151c0317e30b80c33ae")
6 }
7 >

```

(the above code snippet online) because there is no "`Peter Gorilla`" combination in the database. However, it does not allow the following creation:

```

1 > db.employees.insertOne({first_name: "Peter", last_name: "Pan", identity_number:\n
2   "37281923", salary: 3500, contact_details: {}})\n
3 2017-08-10T16:37:00.601+0300 E QUERY      [thread1] WriteError: E11000 duplicate ke\
4 y error collection: accounting.employees index: first_name_1_last_name_1 dup key:\\
5   { : "Peter", : "Pan" } :\n
6 WriteError(\n
7     "index" : 0,\n
8     "code" : 11000,\n
9     "errmsg" : "E11000 duplicate key error collection: accounting.employees i\
10 ndex: first_name_1_last_name_1 dup key: { : \"Peter\", : \"Pan\" }",\n
11     "op" : {\n
12         "_id" : ObjectId("598c617cc0317e30b80c33af"),\n
13         "first_name" : "Peter",\n
14         "last_name" : "Pan",\n
15         "identity_number" : "37281923",\n
16         "salary" : 3500,\n
17         "contact_details" : {\n
18             }\n
19         }\n
20     }\n
21 })\n
22 WriteError@src/mongo/shell/bulk_api.js:469:48\n
23 Bulk/mergeBatchResults@src/mongo/shell/bulk_api.js:836:49\n
24 Bulk/executeBatch@src/mongo/shell/bulk_api.js:906:13\n
25 Bulk/this.execute@src/mongo/shell/bulk_api.js:1150:21\n
26 DBCollection.prototype.insertOne@src/mongo/shell/crud_api.js:242:9\n
27 @(shell):1:1\n
28 >

```

(the above code snippet online)

Because there is already a "Peter Pan" in the database.

Closing Note On Indexes

Building and managing the correct indexes is very important aspect of your MongoDB database design. You might want to read all the details that are given in the official [MongoDB documentation about the indexes](#)

Ruby Integration

So far so good. We have worked with mongo shell, but how can we integrate our Ruby application with MongoDB? Can we build Ruby applications that save data in MongoDB?

We can definitely do. And the answer is [Mongo Ruby Driver](#).

A Ruby MongoDB Application

Let's implement, very quickly, a Ruby MongoDB Application. A simple bookstore. The application is going to be a command line application.

Here is a short demo of this application running:

Mongo Ruby Integration - Command Line App Demo

The text that follows explains all the necessary steps to implement this small application

Project Root Folder

Create the folder that is going to host the source code of your application. Let's name it `bookstore-mongo-ruby-driver`.

```
1 $ mkdir bookstore-mongo-ruby-driver
2 $ cd bookstore-mongo-ruby-driver
3 bookstore-mongo-ruby-driver $
```

(the above code snippet online)

Rvm Integration

Make sure you create the necessary `.ruby-version` and `.ruby-gemset` files so that you integrate your project with Rvm. The `rvm current` command should be returning back that you work with a Ruby 2.X version in a gemset dedicated to your new project.

Something like this:

```
1 bookstore-mongo-ruby-driver $ rvm current
2 ruby-2.2.3@bookstore-mongo-ruby-driver
3 bookstore-mongo-ruby-driver $
```

(the above code snippet online)

Gemfile

Let's go ahead and create our `Gemfile` that will reference the gem necessary to integrate Ruby with MongoDB:

```
1 # File: Gemfile
2 #
3 source 'https://rubygems.org'
4
5 gem 'mongo'
```

(the above code snippet online)

Then run the `bundle` command to bring all the necessary gems for your Ruby / Mongo app to work.

The Main Commands

Before we go into the implementation of the application, let's list very quickly the main Ruby methods that we will use:

1. `require 'mongo'`: This is necessary to bring the Ruby MongoDB integration API.
2. `Mongo::Client.new()`: It will be used to get an instance to a MongoDB client object.
3. `<client_reference> [<collection_name>]`: It will give us access to a specific collection.
4. `<collection_reference>.insert_one()`: It will be used to insert one new document.
5. `<collection_reference>.find()`: It will be used to query a collection.
6. `<collection_reference>.delete_one()`: It will be used to delete a document from a collection.
7. `<collection_reference>.update_one()`: It will be used to update a document.

The `main.rb` Program

Having said the above, this is the code that we have written in order to implement the application:

```

1 require 'mongo'
2
3 $mongo_client = Mongo::Client.new(['127.0.0.1:27017'], database: 'book_store')
4
5 def print_commands
6   print '(1)list books, (i)nsert one, (d)elete, (u)pdate, (q)uit: '
7 end
8
9 def print_book_list
10   $mongo_client[:books].find.each do |book|
11     puts book.inspect
12   end
13 end
14
15 def insert_one
16   print 'Give the title: '
17   title = gets.chomp
18   if title.empty?
19     puts 'Title cannot be empty'
20     return
21   end
22
23   print 'Give the author: '
24   author = gets.chomp
25   if author.empty?
26     puts 'Author cannot be empty'
27     return
28   end

```

```
29
30     print 'Give the tags in comma separated list: '
31     tags = gets.chomp.split(',').map(&:strip)
32
33     book = {title: title, author: author, tags: tags}
34     result = $mongo_client[:books].insert_one(book)
35     puts "Insertion result: #{result.inspect}"
36 end
37
38 def delete_one
39     print 'Give me the title of book to delete: '
40     title = gets.chomp
41     book = $mongo_client[:books].find({title: title}).first
42     unless book
43         puts 'Book not found'
44         return
45     end
46     result = $mongo_client[:books].delete_one({title: title})
47     puts "Deletion result: #{result.inspect}"
48 end
49
50 def update_one
51     print 'Give me the title of the book you want to update: '
52     old_title = gets.chomp
53     book = $mongo_client[:books].find(title: old_title).first
54     unless book
55         puts 'Book not found'
56         return
57     end
58
59     book = {}
60
61     print "Give the new title. Leave empty if you don't want to update: "
62     title = gets.chomp
63     book.merge!({title: title}) unless title.empty?
64
65     print "Give the new author. Leave empty if you don't want to update: "
66     author = gets.chomp
67     book.merge!({author: author}) unless author.empty?
68
69     print "Give the tags in comma separated list. Leave empty if you don't want to \
70 update: "
71     tags = gets.chomp.split(',').map(&:strip)
72     book.merge!({tags: tags}) unless tags.empty?
73
74     result = $mongo_client[:books].update_one({title: old_title}, {"$set" => book})
```

```

75   puts "Update result: #{result.inspect}"
76 end
77
78 print_commands
79
80 while (answer = gets.chomp) != 'q'
81 case
82   when answer == 'l'
83     print_book_list
84   when answer == 'i'
85     insert_one
86   when answer == 'd'
87     delete_one
88   when answer == 'u'
89     update_one
90   when answer == 'q'
91     puts 'Bye!'
92     exit
93 else
94   puts "Wrong Choice! Try again."
95 end
96
97 print_commands
98 end

```

(the above code snippet online)

Let's see the important lines of this program one-by-one:

1. On line 1, we require the `mongo` library file.
2. On line 3, we create a MongoDB client object, using the `Mongo::Client.new()` initilizer. This method takes as first argument an array with the MongoDB server IPS & Ports. You may have a cluster of MongoDB server that is why you are requested to give the IP & Port of your server inside an Array. If you have only one server, the Array should only have one entry. Since we are connecting to the local MongoDB server, we give as IP `127.0.0.1` and as Port `27017`. Then we give a hash of options, the minimum of which should be the database we want to connect to. It is our `book_store` database.
3. The method `print_book_list` shows how we can access the `books` collection of our database. We give the name of the collection as a symbol to the client reference using the square brackets operator `$mongo_client[:books]`.
4. Whenever we have access to the collection we are interested in, then we can execute MongoDB calls, like on line 10, where we call `.find` and we get back an iterator that we can use to go from one result to the next.
5. On line 34, we use `$mongo_client[:books].insert_one({...})` method to create a new book inside the `books` collection. Here, we just give the Hash object like we did with the `insertOne()` when we used the mongo shell JavaScript API.

6. On line 41, we locate a specific document with the help of `.find()` call on the collection reference. Also, we call `.first` in order to get the first element in the collection.
7. Whereas on line 46, we call `$mongo_client[:books].delete_one({title: title})` that will delete the book that has the provided `title`. The `delete_one()` method call takes same arguments as the corresponding `deleteOne()` JavaScript call that we learned in the mongo shell API.
8. On line 73, we call `$mongo_client[:books].update_one({title: old_title}, {"$set" => book})`. This is a call to `update_one()` method which takes the same arguments as the `updateOne()` of the JavaScript API.

That's it. That was a short demo of the MongoDB Ruby driver API. Mostly, it resembles the JavaScript API that we use in the mongo shell.

Ruby on Rails Integration

Before we finish our MongoDB encounter, we will have a look on the MongoDB Ruby on Rails integration.

Obviously, one can use the `mongo` gem, i.e. the MongoDB Ruby driver to integrate with Ruby on Rails, but, there is another gem that has been built on top of that and it provides an ActiveRecord-like interface that proves to be very useful to map the Ruby object world to the document world. This gem is called [mongoid](#).

Create a New Rvm Gemset

Let's start by creating and using new rvm gemset. For example:

```
1 $ rvm use 2.4.0@bookstore-mongoid --create
```

(the above code snippet online)

Install bundler and rails

On this gemset, install `bundler` and then `rails`:

```
1 $ gem install bundler --no-ri --no-rdoc
2 ...
3 $ gem install rails --no-ri --no-rdoc
4 ...
5 $
```

(the above code snippet online)

Initialize The Application

Let's initialize a new bookstore application using `rails`:

```
1 $ rails new bookstore --skip-active-record
2 ...
3 $
```

(the above code snippet online)

The above will make sure that your Ruby on Rails application is initialized without ActiveRecord support. We don't want ActiveRecord because we will be using mongoid instead.

Change directory to the newly created folder bookstore.

Rvm setup

Make sure that you create the .ruby-version and .ruby-gemset files in order to freeze the Ruby version and gemset for Rvm.

Reference haml-rails Gem

Let's integrate with HAML by adding the haml-rails gem into our Gemfile. Then, run bundle to bring all the dependencies in.

Convert Existing ERB files to HAML

Now, let's convert existing ERB files to HAML:

```
1 bookstore $ bin/rake haml:erb2haml
2 -----
3 Generating HAML for app/views/layouts/application.html.erb...
4 Generating HAML for app/views/layouts/mailers.html.erb...
5 Generating HAML for app/views/layouts/mailers.text.erb...
6 -----
7 HAML generated for the following files:
8     app/views/layouts/application.html.erb
9     app/views/layouts/mailers.html.erb
10    app/views/layouts/mailers.text.erb
11 -----
12 Would you like to delete the original .erb files? (This is not recommended unless\
13 you are under version control.) (y/n)
14 y
15 Deleting original .erb files.
16 -----
17 Task complete!
18 No .erb files found. Task will now exit.
19 bookstore $
```

(the above code snippet online)

Reference mongoid Gem

You need to edit your `Gemfile` to reference the `mongoid` gem:

Then run `bundle` to bring `mongoid` in. This will also bring in `mongo` gem, since `mongoid` relies on it.

Mongoid Configuration File

Let's now generate the Mongoid configuration file:

```
1 bookstore $ rails generate mongoid:config
2       create config/mongoid.yml
3 bookstore $
```

(the above code snippet online)

This is what has been generated by default:

```
1 # File: config/mongoid.yml
2 #
3 development:
4   # Configure available database clients. (required)
5   clients:
6     # Defines the default client. (required)
7     default:
8       # Defines the name of the default database that Mongoid can connect to.
9       # (required).
10      database: book_store
11      # Provides the hosts the default client can connect to. Must be an array
12      # of host:port pairs. (required)
13      hosts:
14        - localhost:27017
15      options:
16        # Change the default write concern. (default = { w: 1 })
17        # write:
18        #   w: 1
19
20        # Change the default read preference. Valid options for mode are: :second\
21 ary,
22        # :secondary_preferred, :primary, :primary_preferred, :nearest
23        # (default: primary)
24        # read:
25        #   mode: :secondary_preferred
26        #   tag_sets:
27        #     - use: web
28
```

```
29      # The name of the user for authentication.  
30      # user: 'user'  
31  
32      # The password of the user for authentication.  
33      # password: 'password'  
34  
35      # The user's database roles.  
36      # roles:  
37      #   - 'dbOwner'  
38  
39      # Change the default authentication mechanism. Valid options are: :scram,  
40      # :mongodb_cr, :mongodb_x509, and :plain. Note that all authentication  
41      # mechanisms require username and password, with the exception of :mongod\  
42 b_x509.  
43      # Default on mongoDB 3.0 is :scram, default on 2.4 and 2.6 is :plain.  
44      # auth_mech: :scram  
45  
46      # The database or source to authenticate the user against.  
47      # (default: the database specified above or admin)  
48      # auth_source: admin  
49  
50      # Force a the driver cluster to behave in a certain manner instead of aut\  
51 o-  
52      # discovering. Can be one of: :direct, :replica_set, :sharded. Set to :di\  
53 rect  
54      # when connecting to hidden members of a replica set.  
55      # connect: :direct  
56  
57      # Changes the default time in seconds the server monitors refresh their s\  
58 tatus  
59      # via ismaster commands. (default: 10)  
60      # heartbeat_frequency: 10  
61  
62      # The time in seconds for selecting servers for a near read preference. (\  
63 default: 0.015)  
64      # local_threshold: 0.015  
65  
66      # The timeout in seconds for selecting a server for an operation. (defaul\  
67 t: 30)  
68      # server_selection_timeout: 30  
69  
70      # The maximum number of connections in the connection pool. (default: 5)  
71      # max_pool_size: 5  
72  
73      # The minimum number of connections in the connection pool. (default: 1)  
74      # min_pool_size: 1
```

```
75      # The time to wait, in seconds, in the connection pool for a connection
76      # to be checked in before timing out. (default: 5)
77      # wait_queue_timeout: 5
78
79
80      # The time to wait to establish a connection before timing out, in second\
81 s.
82      # (default: 5)
83      # connect_timeout: 5
84
85      # The timeout to wait to execute operations on a socket before raising an\
86 error.
87      # (default: 5)
88      # socket_timeout: 5
89
90      # The name of the replica set to connect to. Servers provided as seeds th\
91 at do
92      # not belong to this replica set will be ignored.
93      # replica_set: name
94
95      # Whether to connect to the servers via ssl. (default: false)
96      # ssl: true
97
98      # The certificate file used to identify the connection against MongoDB.
99      # ssl_cert: /path/to/my.cert
100
101     # The private keyfile used to identify the connection against MongoDB.
102     # Note that even if the key is stored in the same file as the certificate,
103     # both need to be explicitly specified.
104     # ssl_key: /path/to/my.key
105
106     # A passphrase for the private key.
107     # ssl_key_pass_phrase: password
108
109     # Whether or not to do peer certification validation. (default: true)
110     # ssl_verify: true
111
112     # The file containing a set of concatenated certification authority certi\
113 fications
114     # used to validate certs passed from the other end of the connection.
115     # ssl_ca_cert: /path/to/ca.cert
116
117
118     # Configure Mongoid specific options. (optional)
119     options:
120         # Includes the root model name in json serialization. (default: false)
```

```
121  # include_root_in_json: false
122
123  # Include the _type field in serialization. (default: false)
124  # include_type_for_serialization: false
125
126  # Preload all models in development, needed when models use
127  # inheritance. (default: false)
128  # preload_models: false
129
130  # Raise an error when performing a #find and the document is not found.
131  # (default: true)
132  # raise_not_found_error: true
133
134  # Raise an error when defining a scope with the same name as an
135  # existing method. (default: false)
136  # scope_overwrite_exception: false
137
138  # Raise an error when defining a field with the same name as an
139  # existing method. (default: false)
140  # duplicate_fields_exception: false
141
142  # Use Active Support's time zone in conversions. (default: true)
143  # use_activesupport_time_zone: true
144
145  # Ensure all times are UTC in the app side. (default: false)
146  # use_utc: false
147
148  # Set the Mongoid and Ruby driver log levels when not in a Rails
149  # environment. The Mongoid logger will be set to the Rails logger
150  # otherwise. (default: :info)
151  # log_level: :info
152
153  # Control whether `belongs_to` association is required. By default
154  # `belongs_to` will trigger a validation error if the association
155  # is not present. (default: true)
156  # belongs_to_required_by_default: true
157
158  # Application name that is printed to the mongodb logs upon establishing a
159  # connection in server versions >= 3.4. Note that the name cannot exceed 128 \
160 bytes.
161  # app_name: MyApplicationName
162 test:
163   clients:
164     default:
165       database: bookstore_test
166       hosts:
```

```

167     - localhost:27017
168     options:
169       read:
170         mode: :primary
171         max_pool_size: 1

```

(the above code snippet online)

This is a long configuration file, but most of its lines are commented out. But, it is ready to be used as is. Only make sure that you set the correct name for your development database, on line 10. I have now set it to book_store in order to match the database we worked with earlier in this chapter.

Otherwise, the development and test environment have the correct default values for you to start doing development. When you go to production, you will have to create the corresponding section or rely on your hosting platform settings.

Information: On Heroku, you can integrate with MongoDB via Mongoid if you add the correct add-on. For example, look [here](#) and [here]<https://devcenter.heroku.com/articles/mongolab>).

Book Model

Let's now implement our Book model and integrate it with MongoDB:

```

1 # File: app/models/book.rb
2 #
3 class Book
4   include Mongoid::Document
5
6   field :title
7   field :author
8   field :tags, type: Array
9
10  def tags_joined
11    tags.join(',') rescue ''
12  end
13
14  def tags=(value)
15    write_attribute(:tags, value.split(',').map(&:strip))
16  end
17 end

```

(the above code snippet online)

These are the things that you need to be aware of from the above code:

1. All the models that persist in our MongoDB database as documents in collection, need to include the module `Mongoid::Document`. Line 4.

2. Then we define the fields that constitute the document so that we can have an API to read and write the particular document. Lines 6 - 8.
3. I have defined the method `tags_joined` in order to get the tags into a comma separated list. Lines 10 - 12.
4. I have overridden `tags` assignment so that it can get a comma separated list and write the attribute to the database after splitting it. Lines 14 - 16.

Rails Console Interaction

Let's confirm the integration of our application with the MongoDB server, by starting a Rails console and fetching some books:

```
1 bookstore $ bin/rails c
2 2.4.0 :001 > Book.count
3 => 6
```

(the above code snippet online)

The `Book.count` will count all the book documents in our database.

On the other hand, the `Book.all` will open an iterator for us to browse all the books:

```
1 2.4.0 :004 > Book.all.each { |b| puts b.inspect }
2 #<Book _id: 598b5ee4c0317e30b80c33a9, title: "Of Mice And Men", author: "John Ste\inbeck", tags: ["novel", "classic"]>
3 #<Book _id: 598b5f20c0317e30b80c33aa, title: "The Idiot", author: "Fyodor Dostoyevsky", tags: ["novel", "classic"]>
4 #<Book _id: 598b5f5bc0317e30b80c33ab, title: "Life of Pi", author: "Yann Martel",\tags: ["novel", "modern"]>
5 #<Book _id: 598b5fa5c0317e30b80c33ac, title: "Portable Cosmos", author: "Jones Alexander", tags: ["history", "modern"]>
6 #<Book _id: 598d44ae0ad9864bb54453e9, title: "1984", author: "George Orwell", tag\s: ["novel"]>
7 #<Book _id: 598d44c50ad9864bb54453ea, title: "MongoDB Server", author: "Panayotis Matsinopoulos", tags: ["it", "databases"]>
8 => #<Mongoid::Contextual::Mongo:0x007f9a849e59c8 @cache=nil, @klass=Book, @cri\teria=#<Mongoid::Criteria
9   selector: {}
10  options: {}
11  class: Book
12  embedded: false>
13 , @collection=#<Mongo::Collection:0x70150813329340 namespace=book_store.books>, @\view=#<Mongo::Collection::View:0x70150813329280 namespace='book_store.books' @fil\ter={} @options={}, @cache_loaded=true>
```

(the above code snippet online)

Rails CRUD for Books

Let's now implement the book management part of our bookstore application.

The Routes

We first add the routes for the book resource:

```
1 # File: config/routes.rb
2 #
3 Rails.application.routes.draw do
4   # For details on the DSL available within this file, see http://guides.rubyonrails.org/
5   resources :books
6 end
```

(the above code snippet online)

You can then print the routes as follows:

```
1 bookstore $ bin/rake routes
2      Prefix Verb    URI Pattern          Controller#Action
3       books GET     /books(.:format)      books#index
4             POST    /books(.:format)      books#create
5 new_book GET    /books/new(.:format)    books#new
6 edit_book GET    /books/:id/edit(.:format) books#edit
7    book GET     /books/:id(.:format)    books#show
8           PATCH   /books/:id(.:format)    books#update
9           PUT     /books/:id(.:format)    books#update
10          DELETE  /books/:id(.:format)    books#destroy
11 bookstore
```

(the above code snippet online)

The Controller Actions

This is the `books_controller.rb`:

```
1 # File: app/controllers/books_controller.rb
2 #
3 class BooksController < ApplicationController
4   before_action :find_book, only: [:edit, :update, :show, :destroy]
5
6   def new
7     @book = Book.new
8   end
9
10  def create
11    @book = Book.new(book_params)
12    if @book.save
13      redirect_to edit_book_url(@book)
14    else
15      render :new
16    end
17  end
18
19  def update
20    if @book.update(book_params)
21      redirect_to edit_book_url(@book)
22    else
23      render :edit
24    end
25  end
26
27  def index
28    @books = Book.all.sort({title: 1})
29  end
30
31  def show
32    redirect_to edit_book_url(@book)
33  end
34
35  def destroy
36    @book.destroy
37    redirect_to books_url
38  end
39
40  private
41
42  def find_book
43    @book = Book.find(params[:id])
44  end
45
46  def book_params
```

```

47     result = params.require(:book).permit(:title, :author, :tags_joined)
48     tags_joined = result.delete(:tags_joined)
49     result.merge({ tags: tags_joined })
50   end
51 end

```

(the above code snippet online)

I don't believe there is something that you don't already know here. Look how we use the class Book in order to build and save instances.

1. On line 11, the Book.new(book_params) instantiates a new Book instance.
2. On line 20, the @book.update(book_params) is used to update a specific book.
3. On line 27, the Book.all.sort({title: 1}) is used to fetch books in ascending sorted order.
4. On line 36, the @book.destroy deletes a book from the collection.
5. On line 43, the Book.find(params[:id]) finds the book given its id.

Again, you will find a lot of similarity between Mongoid and ActiveRecord.

However, you need to take into account that Mongoid maps Objects to Documents, whereas ActiveRecord maps Objects to Relations/Tables. Also, ActiveRecord assumes an RDBMS which offers referential integrity and transactions. MongoDB does not offer that.

The Views

The views to support the CRUD actions for books are pretty much standard:

new View

```

1  -# File: app/views/books/new.html.haml
2  -#
3  %h1 Create a New Book
4
5  = render partial: 'books/form'

```

(the above code snippet online)

edit View

```

1  -# File: app/views/books/edit.html.haml
2  -#
3  %h1 Edit Book Details
4
5  = render partial: 'books/form'

```

(the above code snippet online)

_form Partial

Both new and edit view use the _form partial:

```

1  #- File: app/views/books/_form.html.haml
2  -#
3  = form_for @book do |f|
4    = f.label :title
5    = f.text_field :title, placeholder: 'book title'
6    %br
7    = f.label :author
8    = f.text_field :author, placeholder: 'author of book'
9    %br
10   = f.label :tags
11   = f.text_field :tags_joined, placeholder: 'comma separated list of tags'
12   %br
13   = f.submit

```

(the above code snippet online)

index View

And finally, the index view:

```

1  #- File: app/views/books/index.html.haml
2  -#
3  %h1 List Of Books
4  - unless @books.empty?
5    %table
6      %thead
7        %tr
8          %th Title
9          %th Author
10         %th Tags
11         %th
12    %tbody
13    - @books.each do |book|
14      %tr
15        %td= book.title
16        %td= book.author
17        %td= book.tags_joined
18        %td
19          = link_to 'Edit', edit_book_url(book)
20          &nbsp;
21          = link_to 'Delete', book_url(book), method: :delete, data: {confirm: \
22 'Are you sure?'}
```

(the above code snippet online)

Associations

Like ActiveRecord does, Mongoid offers the ability to define associations between documents, and, definitely, this is a big addition to what the simple mongo gem can do.

The associations that Mongoid lets you define are:

1. `embeds_one`
2. `embeds_many`
3. `has_one`
4. `has_many`
5. `has_and_belongs_to_many`
6. `belongs_to`

The first two are used to define nested documents. Either a single nested document (`embeds_one`) or an array of documents.

Example of `embeds_many`

Let's suppose that we want to support comments for each book. We are going to support it as an embedded document.

Comment Model

First, let's create the Comment model:

```

1 # File: app/models/comment.rb
2 #
3 class Comment
4   include Mongoid::Document
5   include Mongoid::Timestamps::Created
6
7   field :text
8
9   embedded_in :book
10 end

```

(the above code snippet online)

Things that you need to be aware of:

1. The model needs to include the module `Mongoid::Document`, even if it is an embedded document.
2. The model of an embedded document is a class that has a reference back to its hosting document. See line 9. We have used the `embedded_in :book` to tell that the document is to be embedded inside a Book document.
3. The Comment model has the field `text` and a `created_at` timestamp which is created with the inclusion of the `Mongoid::Timestamps::Created` module. If you want both `created_at` and `updated_at` you can use the module `Mongoid::Timestamps` instead. The `created_at` timestamp will be set automatically upon creation.

Book Model

The Book Model needs to be enhanced in order to indicate that it embeds many comments:

```

1 # File: app/models/book.rb
2 #
3 class Book
4   include Mongoid::Document
5   ...
6   field :tags, type: Array
7   embeds_many :comments
8
9   def tags_joined
10  ...
11 end

```

(the above code snippet online)

The only difference is line 9, `embeds_many :comments`.

Rails Console Demo

With the above in place, let's try to use this association from the rails console first:

```

1 bookstore $ bin/rails c
2 2.4.0 :001 > b = Book.last
3 => #<Book _id: 598f0e090ad986e2b60082b9, title: "C from theory to practice", aut\
4 hor: "George S. Tselikis; Nikolaos D. Tselikas", tags: ["c", "programming"]>
5 2.4.0 :002 > b.comments << Comment.new({text: 'very useful book, even a little bi\
6 t advanced for me'})
7 => [#<Comment _id: 598fd9ef0ad986fa92abf393, created_at: 2017-08-13 04:47:43 UTC\
8 , text: "very useful book, even a little bit advanced for me">]
9 2.4.0 :003 >

```

(the above code snippet online)

The 002 command above added an embedded document for a new Comment for the Book at hand.

And this is how this book looks like if we fetch it from the database using mongo shell:

```

1 > use book_store
2 switched to db book_store
3 > db.books.find({title: "C from theory to practice"})
4 { "_id" : ObjectId("598f0e090ad986e2b60082b9"), "title" : "C from theory to pract\
5 ice", "author" : "George S. Tselikis; Nikolaos D. Tselikas", "tags" : [ "c", "pro\
6 gramming" ], "comments" : [ { "_id" : ObjectId("598fd9ef0ad986fa92abf393"), "text\
7 " : "very useful book, even a little bit advanced for me", "created_at" : ISODate\
8 ("2017-08-13T04:47:43.643Z") } ] }
9 >

```

(the above code snippet online)

Do you see the array of embedded comments? It has one element inside. Do you also see the timestamp created_at. This has been automatically set.

Comments in our Rails Application

Can we integrate this new feature at the UI level in our bookstore application?

Display Comments

We now introduce the show view as follows:

```

1 -# File: app/views/show.html.haml
2 -#
3 %h1 Book: #{@book.title}
4 %p
5   Author: #{@book.author}
6 %p
7   Tags: #{@book.tags_joined}
8 %h2 Comments
9 %ul#comments-list
10  - @book.comments.sort { |a, b| b.created_at <= a.created_at }.each do |comment|
11    = render partial: 'books/comment', locals: {comment: comment}

```

(the above code snippet online)

This displays the details of a book and its comments at the bottom. But it relies on the books/_comment.html.haml partial which is this:

```

1 -# File: app/views/books/_comment.html.haml
2 -#
3 %li
4   = comment.text
5   = "(#{comment.created_at})"

```

(the above code snippet online)

In order to make this page accessible, you will have to do

1. a change in your `books_controller.rb` so that the `show` action renders the `show` view.
2. a change in your `app/views/books/index.html.haml` so that the title of a book is displayed as a link to corresponding show page.

The above two changes are left to you as an exercise.

Create a Comment

On the show page, we are going to add something more so that the user is able to add a new comment.

<u>New Route</u>

First, we will need to set the correct endpoint. This is the change that we have to do inside our `config/routes.rb` file:

```

1 # File: config/routes.rb
2 #
3 Rails.application.routes.draw do
4   # For details on the DSL available within this file, see http://guides.rubyonrails.org/routing.html
5   resources :books do
6     member do
7       post '/comments', to: 'comments#create'
8     end
9   end
10 end
11 end

```

(the above code snippet online)

Above, we have added a nested resource. The comments inside the books. In fact, we have limited the nested comments routes only for specific books. This is what the `member do` block does. Let's see the routes generated:

```

1 bookstore $ bin/rake routes
2      Prefix Verb    URI Pattern          Controller#Action
3 comments_book POST   /books/:id/comments(.:format)  comments#create
4      books GET     /books(.:format)        books#index
5              POST    /books(.:format)        books#create
6 new_book GET    /books/new(.:format)      books#new
7 edit_book GET    /books/:id/edit(.:format)  books#edit
8      book GET     /books/:id(.:format)      books#show
9              PATCH   /books/:id(.:format)      books#update
10             PUT    /books/:id(.:format)      books#update
11             DELETE  /books/:id(.:format)      books#destroy
12 bookstore $

```

(the above code snippet online)

Do you see the `comments_book` entry? This is the new entry that has been added. The URI Pattern is `/books/:id/comments` and the controller that will handle this is `CommentsController`, on its action `#create`.

Form To Submit New Comment

Let's enhance the `app/views/books/show.html.haml` to have a form for new comment submission:

```

1  -# File: app/views/show.html.haml
2  -#
3  %h1 Book: #{@book.title}
4  %p
5    Author: #{@book.author}
6  %p
7    Tags: #{@book.tags_joined}
8  %h2 Comments
9  %ul#comments-list
10   - @book.comments.sort { |a, b| b.created_at <= a.created_at }.each do |comment|
11     = render partial: 'books/comment', locals: {comment: comment}
12
13 %h3 New Comment
14 = form_tag comments_book_path(@book), id: 'new_comment_form' do
15   = text_field_tag :new_comment
16   = submit_tag 'Submit'

```

(the above code snippet online)

The new lines that we have added are lines 13 to 16. Although this form can submit to `comments#create` action synchronously, we will do an AJAX call instead. For that reason:

Integrate jQuery

You need to add `jquery-rails` in your `Gemfile` and then `bundle`. This will bring in `jQuery` libraries.

Also, amend your `app/assets/javascript/application.js` so that it requires `jQuery`:

```

1 // This is a manifest file that'll be compiled into application.js, which will in\
2 clude all the files
3 // listed below.
4 //
5 // Any JavaScript/Coffee file within this directory, lib/assets/javascripts, or a\
6 ny plugin's
7 // vendor/assets/javascripts directory can be referenced here using a relative pa\
8 th.
9 //
10 // It's not advisable to add code directly here, but if you do, it'll appear at t\

```

```

11 he bottom of the
12 // compiled file. JavaScript code in this file should be added after the last req\
13 uire_* statement.
14 //
15 // Read Sprockets README (https://github.com/rails/sprockets#sprockets-directives\n
16 ) for details
17 // about supported directives.
18 //
19 //=< require 'jquery'
20 //=< require rails-ujs
21 //=< require_tree .

```

(the above code snippet online)

Submit Form With AJAX Request

Now, let's create the file `app/assets/javascript/new_comment.js` that will submit the form with AJAX request.

```

1 // File: app/assets/javascripts/new_comment.js
2 //
3 $(document).ready(function () {
4     var $newCommentForm = $('#new_comment_form');
5
6     $newCommentForm.find('input[type=submit]').on('click', function (e) {
7         e.preventDefault();
8
9         var $newComment = $newCommentForm.find('#new_comment');
10        var $list = $('#comments-list');
11
12        $.ajax({
13            url: $newCommentForm.attr('action'),
14
15            method: 'POST',
16
17            headers: { 'X-CSRF-Token': $('meta[name="csrf-token"]').attr('content') \
18        }), \
19
20            data: {new_comment: $newComment.val().trim() },
21
22            success: function(data, textStatus, jqXHR) {
23                $list.prepend(data);
24            }
25        })
26    });
27 });

```

(the above code snippet online)

We have already used the above technique in the `JavaScript with Rails Applications` chapter. You may want to pay attention to the success handler. The `data` variable is going to contain whatever the server will return back to the browser. We take it and we put it as first child in the list of comments.

Doing so, we assume that the controller is going to return back some HTML fragment (a `li` to be more specific) that will be added to the existing list of comments.

Controller and Action

Let's create the controller and the action:

```

1 # File: app/controllers/comments_controller.rb
2 #
3 class CommentsController < ApplicationController
4   def create
5     @book = Book.find(params[:id])
6     new_comment_text = params[:new_comment]
7     if new_comment_text.blank?
8       head :unprocessable_entity
9     else
10      @comment = Comment.new(text: new_comment_text)
11      @book.comments << @comment
12      render layout: false
13    end
14  end
15 end

```

(the above code snippet online)

The implementation of the `create` action is pretty simple. Please note:

1. `params[:id]` returns the book the comment is going to be attached to.
2. `params[:new_comment]` returns the comment text to be set for the new comment.
3. If the new comment given is blank, then we don't do anything but we return status code `422`, i.e. Unprocessable Entity.
4. Otherwise, we add the comment and we render the `app/views/comments/create.html.haml` view without layout. We do that because this is an AJAX request and we only return back the necessary HTML part that will be appended to the existing DOM/page.

This is the `app/views/comments/create.html.haml`:

```

1 -# File: app/views/comments/create.html.haml
2 -
3 = render partial: 'books/comment', locals: {comment: @comment}

```

(the above code snippet online)

So, the response is going to be the content of the file `app/views/books/_comment.html.haml` dynamically built for the `@comment` that we have just added.

Ready for Demo

Everything is in place. Restart your server, if you have not done that. This will make sure that you have jQuery in place.

This is a video that demonstrates the usage of AJAX.

Adding Comment to A Book

Example of has_one/has_many/belongs_to

On the other hand, the `has_one/has_many/belongs_to` association is used to reference another external document. We are going to create the model Publisher and then associate book to a Publisher via the `has_many` association.

Publisher Model

Let's suppose that we have the Publisher model as follows:

```
1 # File: app/models/publisher.rb
2 #
3 class Publisher
4   include Mongoid::Document
5
6   field :name
7
8   has_many :books
9 end
```

(the above code snippet online)

The document Publisher is now associated to a collection of books.

Book Model

Let's now amend the Book model accordingly?

```

1 # File: app/models/book.rb
2 #
3 class Book
4   include Mongoid::Document
5
6   field :title
7   field :author
8   field :tags, type: Array
9   embeds_many :comments
10  belongs_to :publisher
11
12  def tags_joined
13    tags.join(',') rescue ''
14  end
15
16  def tags=(value)
17    write_attribute(:tags, value.split(',').map(&:strip))
18  end
19 end

```

(the above code snippet online)

The new line here is the `belongs_to :publisher`.

Rails Console Demo

The following is a rails console session that demonstrates the use of the newly created association:

```

1 2.4.0 :008 > b = Book.where({title: "1984"}).first
2 => #<Book _id: 598d44ae0ad9864bb54453e9, title: "1984", author: "George Orwell", \
3   tags: ["novel", "classic"], publisher_id: nil>
4 2.4.0 :009 > b.publisher
5 => nil
6 2.4.0 :010 > p = Publisher.create({name: 'Harvill Secker'})
7 => #<Publisher _id: 599056fe0ad986042ad001a5, name: "Harvill Secker">
8 2.4.0 :011 > b.publisher = p
9 => #<Publisher _id: 599056fe0ad986042ad001a5, name: "Harvill Secker">
10 2.4.0 :012 > b.save!
11 => true
12 2.4.0 :013 > b.reload
13 => #<Book _id: 598d44ae0ad9864bb54453e9, title: "1984", author: "George Orwell", \
14   tags: ["novel", "classic"], publisher_id: BSON::ObjectId('599056fe0ad986042ad001\
15 a5')>
16 2.4.0 :014 > b.publisher
17 => #<Publisher _id: 599056fe0ad986042ad001a5, name: "Harvill Secker">
18 2.4.0 :015 > p.reload
19 => #<Publisher _id: 599056fe0ad986042ad001a5, name: "Harvill Secker">

```

```

20 2.4.0 :016 > p.books
21 => [#<Book _id: 598d44ae0ad9864bb54453e9, title: "1984", author: "George Orwell"\ 
22 , tags: ["novel", "classic"], publisher_id: BSON::ObjectId('599056fe0ad986042ad00\ 
23 1a5')>]
24 2.4.0 :017 >

```

(the above code snippet online)

1. On command 008 we locate a book. We use the `where()` method and then we take the first that matches.
2. On command 009, we see that this book does not have a Publisher.
3. On command 010, we create a new Publisher.
4. On command 011, we associate the new Publisher to the book and then on command 012 we commit the update to the database.
5. On commands 013 and 014, we make sure that the book has the Publisher as expected.
6. On commands 015 and 016, we see the list of books that this Publisher owns.

mongo Shell Output

Let's see also how the book with title 1984, which belongs to a Publisher, is now displayed when fetched as JSON object from database using JavaScript API:

```

1 > db.books.find({title: "1984"})
2 { "_id" : ObjectId("598d44ae0ad9864bb54453e9"), "title" : "1984", "author" : "Geo\
3 rge Orwell", "tags" : [ "novel", "classic" ], "publisher_id" : ObjectId("599056fe\ 
4 0ad986042ad001a5") }
5 >

```

(the above code snippet online)

Do you see the `publisher_id` property? This is there because of the `belongs_to :publisher` association that we have inserted in the Book model.

And this is the Publisher associated:

```

1 > db.publishers.find("599056fe0ad986042ad001a5")
2 { "_id" : ObjectId("599056fe0ad986042ad001a5"), "name" : "Harvill Secker" }
3 >

```

(the above code snippet online)

Rails App UI

And let's see how we can enhance our bookstore app to support the setting of a Publisher, while creating or editing a Book.

First change: We have to enhance the show view to display the Publisher of a book:

```

1  #- File: app/views/show.html.haml
2  -
3  %h1 Book: #{@book.title}
4  ...
5  %p
6    Publisher: #{@book.publisher.try(:name)}
7  %p
8  ...
9  = submit_tag 'Submit'

```

(the above code snippet online)

We have added lines 6 and 7. We display the Publisher name.

Then we have to enhance the form that allows creation and update of a Book:

```

1  #- File: app/views/books/_form.html.haml
2  -
3  ...
4  %br
5  = f.label :publisher
6  = f.collection_select :publisher_id, Publisher.all, :id, :name, prompt: 'Publis\
7 her?'
8  ...
9  = f.submit

```

(the above code snippet online)

We have added the label and the select box to select the Publisher.

Finally, we have to update the `books_controller.rb` so that it allows the `:publisher_id` attribute.

```

1  # File: app/controllers/books_controller.rb
2  #
3  class BooksController < ApplicationController
4  ...
5  def book_params
6    result = params.require(:book).permit(:title, :author, :tags_joined, :publish\
7 er_id)
8  ...
9  end

```

(the above code snippet online)

We have added the `:publisher_id` in the list of permitted attributes.

Now, everything is ready. You can set the Publisher of a book:

Note: You need to create the Publisher CRUD (Create, Retrieve, Update, Delete) part of the app as an exercise.

Mongoid - Setting Book Publisher

Validations

Mongoid allows us to add validations to our models. This is because it includes the `ActiveModel::Validations` module.

We leave this to you as an exercise:

1. Book validations:
 1. Make `title` mandatory and unique
 2. Make `author` mandatory
 3. Make `publisher` mandatory
2. Publisher validations:
 1. Make `name` mandatory and unique

Closing The MongoDB Chapter

MongoDB is used in many production applications. It is definitely a knowledge that will be proven useful in the job. In order to master it, you need to study very well the [MongoDB documentation](#) itself and then the [Ruby driver](#) and [Mongoid](#) gems documentations.

Tasks and Quizzes

Before you continue, you may want to know that: You can sign up to [Tech Career Booster](#) and have a mentor evaluate your tasks, your quizzes and, generally, your progress in becoming a Web Developer. Or you can sign up and get access to Tech Career Booster Slack channel. In that channel, there are a lot of people that can answer your questions and give you valuable feedback.

Quiz

The quiz for this chapter can be found [here](#)

Task Details

You need to make sure that you carry out whatever the chapter content teaches you. Such as:

1. Using Mongo Ruby driver implement the simple bookstore command line interface application.
2. Using mongoid implement the Rails application that has
 1. All basic CRUD operations for books
 2. Each book should embed many Comments
3. You need to make sure that the books index page can take you to the show page of a book
4. You need to make sure that the controller `show` action renders the `show` view.
5. You need to implement the addition of a comment using AJAX.

6. You need to create all the CRUD pages for the Publisher model.
7. You need to make sure that we can assign a Publisher to a Book
8. You need to implement the validations on Book and Publisher as the chapter enlists them.

Important: Your work needs to be uploaded on your Github account. Also, you need to deploy your application on Heroku and set up integration with production Mongo.